

CS2102

DBMS

Challenges

Want

- (Efficiency) Fast access to information in huge volumes of data
- (Transactions) “All-or-nothing” changes to data
- (Data integrity) Parallel access and changes to data
- (Recovery) Fast and reliable handling of failures
 - HDD/SSD/system crash | Power outage | Network disruption
- (Security) Fine-grained data access rights

File-based data management

- Complex, low-level code
- Often similar requirements across different programs

Problems High development effort | Long development times | Higher risk of (critical) errors

Transaction

- Finite sequence of database operations (reads and/or writes)
- Smallest logical unit of work from an application perspective

Each transaction T satisfies the following ACID properties:

- Atomicity: either all effects of T are reflected in the database or none
- Consistency: Execution of T guarantees to yield a correct state of the database
- Isolation: Execution of T is isolated from the effects of concurrent transactions
- Durability: After the commit of T, its effects are permanent even in case of failures

Concurrent Execution

Common problems

T ₁ (B, 500)	T ₂ (B, 100)
begin	
read(B)	
B = B + 500	
	begin
	read(B)
	B = B + 100
write(B)	
commit	
	write(B)
	commit

Final balance B = 1,100
(effect of T₁, overwritten)

→ Lost Update

T ₁ (B, 100)	T ₂ (B, 500)
begin	
read(B)	
B = B + 500	
write(B)	
	begin
	read(B)
	B = B + 100
	write(B)
	commit
abort	

Final balance B = 1,600
(when it should be 1,100)

→ Dirty Read

T ₁ (B, 100)	T ₂ (B, 500)
begin	
read(B)	
	begin
	read(B)
	B = B + 100
	write(B)
	commit
read(B)	
...	

Balance B is retrieved twice
but the values differ

→ Unrepeatable Read

Requirements

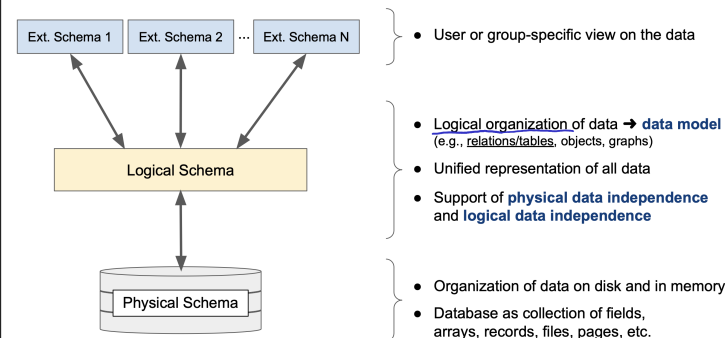
- Want serializable transaction execution
 - A concurrent execution of a set of transactions is **serializable** if this execution is equivalent to some serial execution of the same set of instructions
 - Two executions are equivalent if they have the same effect on the data
- Hence, DBMS
 - Supports concurrent executions of transactions to optimize performance
 - Enforces serializability of concurrent executions to ensure integrity of data

DBMS


- Is a set of universal and powerful functionalities for data management
- Complex, low-level code moved from application logic to DBMS

Benefits Faster application development | Increased productivity | Higher stability / less errors

Data abstraction



Data independence



Logical data independence

- Ability to change logical schema without affecting external schemas (e.g., adding/deleting/updating attributes, changing data types, changing data model)

Physical data independence

- Representation of data independent from physical scheme
- Physical schema can be changed without affecting logical schema (e.g., creating indexes, new caching strategies, different storage devices)

Terminology

Data model

- Set of concepts for describing data
- Framework to specify structure of a DB
- e.g. Relational model, where everything is a table

Schema Description of structure of a DB, using the concepts provided by data model

Schema instance Content of a DB at a particular time

Relational Model

Terminology

Relation

- Set of tuples (or records)
- $R(A_1, A_2, \dots, A_n)$: relation schema with name R , and n attributes A_1, A_2, \dots, A_n
- Each instance of schema R is a relation which is a subset of $\{(a_1, a_2, \dots, a_n) \mid a_i \in \text{dom}(A_i) \cup \{\text{null}\}\}$

Relation schema

- Definition of a relation
- Specifies relation name, attributes (columns) and data constraints (e.g. domain constraints)
 - Employees (id: **integer**, name: **text**, dob: **date**, salary: **numeric**)

Domain Set of atomic values (e.g. integer, numeric, text)

- Domain of attribute A_i , $\text{dom}(A_i)$ = set of possible values of A_i
- Each value v of attribute A_i : either $v \in \text{dom}(A_i)$ or $v = \text{null}$
- **null** - special value indicating the v is not known or not specified

Relational database schema Set of relation schemas + data constraints

Relational database Collection of tables

DB vs DBS vs DBMS

$$DBS = DBMS + n \times DB, \text{ where } n > 0$$

Integrity constraints

- Condition that restricts what constitutes valid data
- 3 main structural integrity constraints of the Relation Model
 - Domain constraints | Key constraints | Foreign key (referential integrity) constraints

Domain constraints

- Data type
- NOT NULL / UNIQUE / PRIMARY KEY / FOREIGN KEY / CHECK / DEFAULT

Key constraints

Superkey Subset of attributes that **uniquely** identifies a tuple in a relation

Key Superkey that is also minimal (i.e. no proper subset of the key is a superkey)

Candidate keys Set of all possible keys for a relation

Primary key

- Chosen candidate key for a relation
- Cannot be **null** (entity integrity constraint)
- Underlined in relation schema
 - Employees (id: **integer**, name: **text**, dob: **date**, salary: **numeric**)

Prime attribute Attribute of a primary key (cannot be null)

Foreign key constraints

Foreign key

- Subset of attributes of relation A that refer to the primary key in a relation B
- Each foreign key in referencing relation must either
 - Appear as primary key in referenced relation, or
 - Be **null**

Properties

- Specified by DB designer to define what constitutes valid data
- Referencing and referenced relation can be the same relation (e.g. each employee has at most one manager)
- Relation can be referencing and referenced relation for different relations

Limitations

- Covers application-independent constraints (e.g. limit domain to valid values)
- Does not cover application-dependent constraints derived from deeper semantics of the data

Practical considerations

- Optional, not mandatory
- May affect performance, since checking constraints require additional processing

RELATIONAL ALGEBRA

- Relations are closed under the Relational Algebra

Unary operators

Selection σ_c

- For each tuple $t \in R, t \in \sigma_c(R) \iff$ selection condition c evaluates to true for tuple t .
- Input and output have same schema
- e.g. Find all projects where Judy is the manager:
 $\sigma_{\text{manager}=\text{'Judy'}}(\text{Projects})$

Selection condition is a boolean expression of one of the following forms:

expression	example
attribute op constant	$\sigma_{\text{start}=2020}(\text{Projects})$
$attr_1$ op $attr_2$	$\sigma_{\text{start}=\text{end}}(\text{Projects})$
$expr_1 \wedge expr_2$	$\sigma_{\text{start}=2020 \wedge \text{end}=2021}(\text{Projects})$
$expr_1 \vee expr_2$	$\sigma_{\text{start}=2020 \vee \text{end}=2021}(\text{Projects})$
$\neg expr$	$\sigma_{\neg(\text{start}=2020)}(\text{Projects})$
$(expr)$	-

where

- op** $\in \{=, <>, <, \leq, \geq, >\}$
 - Precedence: $()$, **op**, \neg , \wedge , \vee
 - Comparision with **null** is **unknown**, arithmetic with **null** is **null**
- In boolean expressions, treat unknown as literally unknown. e.g.
- $\text{false} \wedge \text{unknown} = \text{false}$
 - $\text{false} \vee \text{unknown} = \text{unknown}$
 - $\neg \text{unknown} = \text{unknown}$
 - $\text{true} \wedge \text{unknown} = \text{unknown}$
 - $\text{true} \vee \text{unknown} = \text{true}$

Projection π_l

- Projects columns of a table specified in list l
- Order of attributes in l matters
- Duplicates are removed, because a relation is a set of tuples

Example

Teams			$\pi_{\text{pname,ename}}(\text{Teams})$	
ename	pname	hours	pname	ename
Sarah	BigAI	10	BigAI	Sarah
Sam	BigAI	5	BigAI	Sam
Sam	BigAI	3		

Renaming ρ_l

- Renames attributes of a relation

Consider $R(\text{ename}, \text{pname}, \text{hours})$. Rename ename to name, pname to title. Can either specify

- list of all attr.: $\rho_{(\text{name}, \text{title}, \text{hours})}(R)$
- or list of renames:
 $\rho_{\text{name} \leftarrow \text{ename}, \text{title} \leftarrow \text{pname}}(R)$

Set operations

- Union, Intersection, Set difference (all obvious)
- Note: intersection can be expressed with union and set difference:
 $R \cap S = (R \cup S) - ((R - S) \cup (S - R))$
- The two relations must be union-compatible

Union compatability

Two relations are union-compatible if

- Same number of attributes
- Corresponding attributes have same or compatible domains (different attribute names are ok)

Example The following are union-compatible.

- Employees(name: **text**, role: **text**, age: **integer**)
- Teams(ename: **text**, pname: **text**, hours: **integer**)

Cross product

Forms all possible pairs of tuples from the two relations

Division operator

- This question considers a binary relational algebra operator called the **division operator** denoted by $/$ ¹.
Consider two relations R and S where the set of attributes in the schema of R and S are $(A_1, \dots, A_m, B_1, \dots, B_n)$ and (B_1, \dots, B_n) respectively where $m \geq 1$ and $n \geq 1$. That is, the set of attributes in S is a *proper* subset of the set of attributes in R .
Assume that the attributes that are in R but not in S are ordered as (A_1, \dots, A_m) in the schema of R and the schema of S is (B_1, \dots, B_n) . Let L denote the list of attributes in the schema of R .
The division of R by S (denoted by R/S) computes the largest set of tuples $Q \subseteq \pi_{A_1, \dots, A_m}(R)$ such that for every tuple $(a_1, \dots, a_m) \in Q$,
$$\pi_L(\{(a_1, \dots, a_m)\} \times S) \subseteq R$$

 Q is also referred to as the **quotient** of R/S and its schema is (A_1, \dots, A_m) . The following example illustrates R/S given two relations $R(A, B)$ and $S(B)$.

R		S	R/S
A	B		A
a	1	B 1 2	a c
a	2		
b	1		
c	1		
c	2		
c	3		
d	2		
d	3		

Join operations

- Combines \times, σ_c, π_l into a single op
- Simple relational algebra expressions

Inner joins

- Eliminates tuples that do not satisfy matching criteria (i.e. selection)
- Is a selection from cross product

θ-Join

R ⋈_θ S = σ_θ(R × S)

Equi Join

Like θ-Join, but θ must only involve =

Natural Join

Like equi join (i.e. only equality operator), but

- Join is performed over common attributes of R and S
- If there are no common attributes, acts like a cross product, since selection condition c is vacuously true
- Output relation keeps one copy of common attributes

Formally,

R ⋈ S = π_l(R ⋈_c ρ_{b_i ← a_i, ..., b_k ← a_k}(S))

where

- A = {a_i, ..., a_k} is the set of common attributes of R and S
- c = (a_i = b_i) ∧ ... ∧ (a_k = b_k)
- l = list of (attr. of R + attr. of S not in A)

Outer joins

- Inner join + dangling tuples
- A **dangling tuple** is a tuple that doesn't satisfy the inner join condition, i.e. foreign key not referenced in the relation.

Steps

- Perform inner join M = R ⋈_θ S
- To M, add dangling tuples from

{ R in left outer join ⋈_l
S in right outer join ⋈_r
R and S in full outer join ⋈_f }

- Pad missing attribute values with **null**

Formal definitions

- Set of dangling tuples in R, with respect to R ⋈_θ S
dangle(R ⋈_θ S) ⊆ R
- null(R) is a n-compoennt tuple of **null** values, where n is the number of attributes in R
- Left outer join (R ⋈_l S)
= (R ⋈_θ S) ∪ (dangle(R ⋈_θ S) × {null(S)})
- Right outer join (R ⋈_r S)
= (R ⋈_θ S) ∪ ({null(R)} × dangle(S ⋈_θ R))
- Full outer join (R ⋈_f S)
= (R ⋈_θ S) ∪ ((dangle(R ⋈_θ S) × {null(S)}) ∪ ({null(R)} × dangle(S ⋈_θ R)))

Natural outer joins

- Like natural inner joins
- Only equality operator used for condition
- Join is performed over common atributes of R and S
- Output relation keeps one copy of common attributes

Complex expressions

There are multiple ways to formulate a query to get the same result, e.g.

- Order of joins
- Order of selection (before/after join)
- Additional projections to minimize intermediate results

Invalid expressions

- Attribute no longer available after projection
σ_{role='dev'}(π_{name,age}(Employees))
- Attribute no longer available after renaming
σ_{role='dev'}(ρ_{position ← role}(Employees))
- Incompatible attribute types
σ_{age=role}(Employees)

SQL (DEFINE & MANIPULATE)

- Declarative language: focus on what to compute, not on how to compute
- Statement Level Interface - app is a mixture of host language statements and SQL statements
- Call Level Interface - app is written in host language, SQL statements passed as arguments

Data types

- **CAST(x AS NUMERIC)** for typecasting
 - Useful to force floating point division

type	description
boolean	logical Boolean (true/false)
integer	signed 4-byte integer
float8	double precision floating-point number (8 bytes)
numeric(p, s)	number with p significant digits and s decimal places
char(n)	fixed-length character string
varchar(n)	variable-length character string
text	variable-length character string
date	calendar date (year, month, day)
timestamp	date and time

Other extended types:

- Document types: XML, JSON
- Spatial types: point, line, polygon, circle, box, path
- Special types: money/currency, MAC/IP address
- User defined types

Data definition

Creating tables

Employees (id: **integer**, name: **text**, dob: **date**, salary: **numeric**)

```
CREATE TABLE Employees(  
  id    INTEGER,  
  name  VARCHAR(50),  
  age   INTEGER,  
  role  VARCHAR(50) DEFAULT 'sales'  
);
```

Modifying a schema

```
-- change data type  
ALTER TABLE Projects ALTER COLUMN name TYPE VARCHAR(200);  
  
-- set default value  
ALTER TABLE Projects ALTER COLUMN start_year SET DEFAULT 2021;  
  
-- drop default value  
ALTER TABLE Projects ALTER COLUMN start_year DROP DEFAULT;  
  
-- add new column with default value  
ALTER TABLE Projects ADD COLUMN budget NUMERIC DEFAULT 0.0;  
  
-- drop column from table  
ALTER TABLE Projects DROP COLUMN budget;  
  
-- add FK constraint  
ALTER TABLE Teams ADD CONSTRAINT eid_fkey FOREIGN KEY (eid) REFERENCES Employees(id);  
  
-- drop FK constraint  
ALTER TABLE Teams DROP CONSTRAINT eid_fkey;  
  
-- drop table  
DROP TABLE Projects;  
DROP TABLE IF EXISTS Projects; -- avoids throwing error if does not exist  
  
-- drop table with dependent objects  
DROP TABLE Projects CASCADE; -- deletes Projects and FK constraint
```

Data manipulation

Insert data

```
-- Specify all attribute values
INSERT INTO Employees VALUES (101, 'Sarah', 25, 'dev');
-- Specify selected attribute values
-- role: sales, age: NULL
INSERT INTO Employees (id, name) VALUES (102, 'Judy');
```

Deleting data

```
-- Delete all tuples
DELETE FROM Employees;
-- Delete selected tuples
DELETE FROM Employees WHERE role = 'dev';
```

Updating data

```
-- Update with where clause
UPDATE Employees SET age = age+1 WHERE name='Sarah';
-- Set all age to 0
UPDATE Employees SET age = 0;
-- Uppercase ALL strings
UPDATE Employees SET name=UPPER(name), role=UPPER(role);
```

Handling NULL

- Comparison with **null** is unknown
- Arithmetic opertaion with **null** is **null**

IS (NOT) NULL

- value is **null** \iff evaluates to true
- x IS NOT NULL** equivalent to **NOT (x IS NULL)**

IS (NOT) DISTINCT FROM

- Equivalent to **x <> y** if *x* and *y* both non-null
- Both null \implies return false
- One null \implies return true

Constraints

- All constraints can be named or unnamed (unnamed constraints still get named by DBMS)
- All column constraints can be specified as table constraints (except not null)
- Table constraints referring to a single column can be specified as column constraint
- Column and table constraints can be combined, e.g.

```
CREATE TABLE Employees (
  -- id specified twice
  id INTEGER NOT NULL,
  name VARCHAR(50),
  UNIQUE(id)
);
```

Not-null constraints

- Violation: $\exists t \in \text{Employees}$ where **t.id IS NOT NULL** evaluates to false

```
CREATE TABLE Employees(
  -- unnamed (name assigned by DBMS)
  id INTEGER NOT NULL,
  -- named (easier bookkeeping)
  name VARCHAR(50) CONSTRAINT nn_name NOT NULL
);
```

Unique constraints

- Violation: For any two tuples $x, y \in \text{Teams}$, (**x.id <> y.id**) or (**x.name <> y.name**) evaluates to false
- Since (**null <> null**) evaluates to unknown, the tuples (101, **null**) and (101, **null**) are considered unique
- Unique constraint involving multiple attributes is specified using table constraints

```
CREATE TABLE Employees (
  -- column constraint
  id INTEGER UNIQUE,
  name VARCHAR(50),
  -- table constraint
  UNIQUE(id), -- single attribute
  UNIQUE(id, name), -- multi attribute
  CONSTRAINT unique_id UNIQUE(id) -- named
);
```

PK constraints

- PRIMARY KEY** and **UNIQUE NOT NULL** have the same effect
- Only 1 **PRIMARY KEY**, but can have multiple **UNIQUE NOT NULL**

```
CREATE TABLE Employees (
  -- These 2 statements have the same effect
  id INTEGER PRIMARY KEY,
  id INTEGER UNIQUE NOT NULL,
  --
  name VARCHAR(50),
  PRIMARY KEY (id, name) -- multiple attributes
);
```

FK constraints

The foreign key must either

- exist in the other table, or
- contain NULL in at least one of its attributes

```
CREATE TABLE Employees (
  id INTEGER PRIMARY KEY,
  name VARCHAR(50),
  age INTEGER
);
```

```
CREATE TABLE Projects (
  name VARCHAR(50) PRIMARY KEY,
  start_year INTEGER,
  end_year INTEGER
);
CREATE TABLE Teams (
  eid INTEGER, -- eid -> Employee.name
  pname VARCHAR(100), -- pid -> Project.name
  hours INTEGER,
  PRIMARY KEY (ename, pname),
  FOREIGN KEY (eid) REFERENCES Employees (id),
  FOREIGN KEY (pname) REFERENCES Projects (name)
);
```

FK constraints ON action

- Specify action in case a FK constraint is violated

• **ON DELETE/UPDATE <action>**

- Both specifications are optional

case	action
NO ACTION (default)	rejects delete/update if it violates constraint
RESTRICT	similar to NO ACTION, but check of constraint cannot be deferred
CASCADE	propagates delete/update to referencing tuples (can significantly affect performance)
SET DEFAULT	updates FKs of referencing tuples to some default value (default value must be a PK in referenced table)
SET NULL	updates FKs of referencing tuples to null (corresponding column must be allowed to contain null values)

```
CREATE TABLE Teams (
  eid INTEGER,
  pname VARCHAR(100),
  hours INTEGER,
  PRIMARY KEY (ename, pname),
  FOREIGN KEY (eid) REFERENCES Employees (id) ON DELETE
    <action> ON UPDATE <action>,
  FOREIGN KEY (pname) REFERENCES Projects (name) ON
    DELETE <action> ON UPDATE <action>
);
```

Check constraints

- Specify that column values must satisfy a Boolean expression
- Scope: one table, single row

```
-- column constraint
CREATE TABLE Teams (
  eid    INTEGER PRIMARY KEY,
  pname  VARCHAR(100),
  -- unnamed
  hours  INTEGER check (hours > 0),
  -- named
  hours  INTEGER constraint positive_hours check (hours
    > 0)
);

-- table constraint
CREATE TABLE Projects (
  name      VARCHAR(50) PRIMARY KEY,
  start_year INTEGER,
  end_year  INTEGER,
  CHECK (start_year <= end_year)
);
```

Can be complex Boolean expressions

```
CREATE TABLE Teams (
  eid    INTEGER PRIMARY KEY,
  pname  VARCHAR(100),
  hours  INTEGER,
  CHECK (
    (pname = 'CoreOS' AND hours >= 30)
    OR
    (pname <> 'CoreOS' AND hours >= 0)
  )
);
```

Deferrable constraints

- Default behavior for constraints: check immediately at the end of SQL statement execution, even for transaction with multiple SQL statements
 - Violation causes statement to be rolled back

Deferrable constraints

- Check can be deferred for some constraints to the end of a transaction
- Allows violation of constraints temporarily within the scope of a transaction
- Constraint still needs to be resolved at the end of the transaction
- Can be used for **UNIQUE**, **PRIMARY KEY**, **FOREIGN KEY**

Benefits

- No need to care about order of SQL statements within a transaction

- Allows for cyclic FK constraints
- Performance boost when constraint checks are bottleneck

Downsides

- Difficult to troubleshoot
- Data definition no longer unambiguous
- Performance penalty when performing queries


ER MODEL

Entity

- Objects that are distinguishable from other objects
- Entity set**: Collection of entities of the same type

Attribute

- Specific information describing an entity
- Key attribute(s)** uniquely identifies each entity
- Composite attribute** composed of multiple other attributes
- Multivalued attribute** may consist of more than one value for a given entity
- Derived attribute** derived from other attributes



Relationship

Association among two or more entities

Relationship set

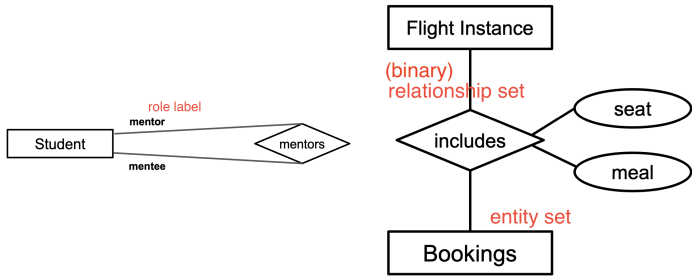
- Collection of relationships of the same type
- Can have their own attributes that further describe the relationship
- $Key(E_i)$ is the attributes of the selected key of entity set E_i

Role

- Describes an entity set’s participation in a relationship
- Explicit role label only in case of ambiguities (e.g. same entity set participates in same relationship more than once)


Degree

- An n -ary relationship set involves n entity roles, where n is the degree of the relationship set
- Typically binary or ternary




Cardinality constraints

- Upper bound** for entity’s participation




Participation constraints

- Lower bound** for entity’s participation
- Partial (default): participation not mandatory
- Total: mandatory (at least 1)



Alternative



Implementation

Many-to-Many Represent relationship set with a table

Many-to-One

- 1. Represent relationship set with a table
- 2. Combine relationship set and total participation entity set into one table

One-to-One

- 1. Represent relationship set with a table
- 2. Combine relationship set and either entity set into one table


Dependency constraints

Weak entity sets

- Entity set that does not have its own key
- Can only be uniquely identified by considering primary key of owner entity
- Existence depends on existence of owner entity

Partial key

- Set of attributes of weak entity set that uniquely identifies a weak entity, for a given owner entity



Requirements

- Many-to-one relationship from weak entity set to owner entity set
- Weak entity set must have total participation in identifying relationship

Relational mapping

- Entity set → table
- Composite/multivalued attributes:
 - 1. Convert to single-valued attributes
 - 2. Additional table with FK constraint
 - 3. Convert to a single-valued attribute (e.g. comma separated string)

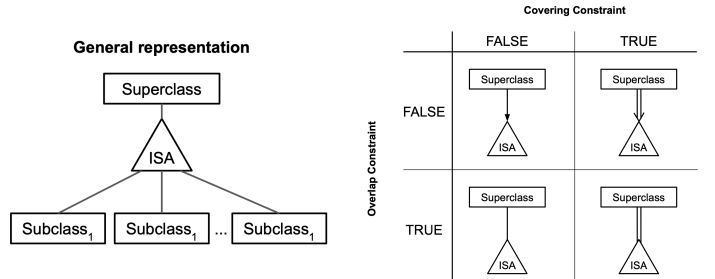
ISA Hierarchies

- “Is a” relationship - used to model generalization/specialization of entity sets

Constraints

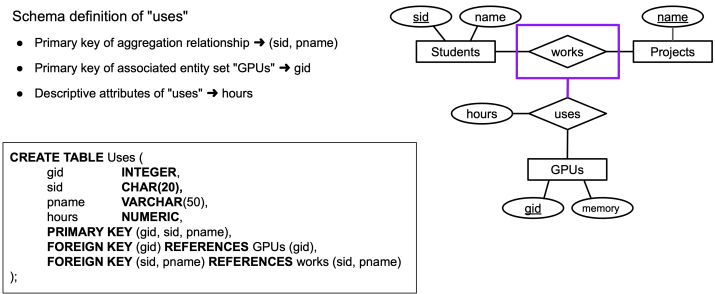
Overlap Can a superclass entity belong to multiple subclasses?

Covering Does a superclass entity have to belong to a subclass?



Aggregation

- Abstraction that treats relationships as higher-level entities



SQL (QUERIES)

Basic form

```
SELECT DISTINCT a1, a2, ... am
FROM r1, r2, ... rm
WHERE c
```

corresponds to

$$\pi_{a_1, a_2, \dots, a_m}(\sigma_c(r_1 \times r_2 \times \dots \times r_n))$$

SELECT

- Wildcard '*' to include all attributes
- expr BETWEEN <lower> AND <upper> for basic value range conditions

```
SELECT * FROM countries
WHERE (continent = 'Asia' OR continent = 'Europe')
AND (population BETWEEN 5000000 AND 6000000)
```

- || for string concatenation

```
-- AS is optional
SELECT name, '$$ ' || ROUND((gdp/population)*1.38)
AS gdp_per_capita FROM countries;
```

- SELECT DISTINCT to eliminate duplicates

- Two tuples (n1, c1) and (n2, c2) are distinct if (n1 IS DISTINCT FROM n2) or (c1 IS DISTINCT FROM c2) evaluates to true

WHERE

- Returns rows that evaluate to true
- Does not return rows that evaluate to unknown/null!
- Use IS (NOT) NULL for comparison with null
- IS (NOT) LIKE
 - ‘_’ matches single char
 - ‘%’ matches any sequence of zero or more chars

Set operations

- Eliminate duplicates: UNION, INTERSECT, EXCEPT
- Keep duplicates: UNION ALL, INTERSECT ALL, EXCEPT ALL

Join operations

- JOIN is interpreted as INNER JOIN
- NATURAL JOIN
- LEFT OUTER JOIN interpreted as LEFT JOIN
 - Use WHERE ... IS NULL to get only dangling tuples

Subqueries

In FROM clause

- Must be enclosed in parenthesis
- Table alias mandatory
- Column aliases optional

```
SELECT * FROM (
  SELECT n.iso2, n.name
  FROM countries n, borders b
  WHERE n.iso2 = b.country1_iso2
  AND country2_iso2 IS NULL
) AS LandborderfreeCountries;
-- column aliases optional
-- ) AS LandborderfreeCountries(code, name);
```

IN subquery

- expr IN (subquery)
- Subquery must return exactly one column
- Returns true if expr matches with any subquery row
- IN can typically be replaced with inner joins
- NOT IN can typically be replaced with outer joins

```
-- subquery is a SELECT
SELECT * FROM countries WHERE name IN (SELECT name FROM cities);
-- subquery is a manually specified result of a subquery
SELECT * FROM countries WHERE continent IN ('Asia', 'Europe');
```

ANY/SOME subquery

- `expr op ANY (subquery)`
- Subquery must return exactly one column
- Expression is compared to each subquery row using operator `op`
- Returns true if comparison evaluates to TRUE for at least one subquery row
- `.. < ANY(..) ⇒` Not maximum
- `.. > ANY(..) ⇒` Not minimum

ALL subquery

- `expr op ANY (subquery)`
- Subquery must return exactly one column
- Expression is compared to each subquery row using operator `op`
- Returns true if comparison evaluates to TRUE for all subquery rows
- `.. <= ALL(..) ⇒` Minimum
- `.. >= ALL(..) ⇒` Maximum

EXISTS subqueries

- `EXISTS (subquery)`
- Returns TRUE if subquery returns at least one tuple
- Generally correlated. If uncorrelated, then likely either wrong or unnecessary

Correlated subqueries

- Uses value from outer query
- Result of subquery depends on value of outer query
 - Potentially slow performance
 - For `ALL` condition, problematic if subquery contains NULL value, since condition never evaluates to TRUE
 - Potential naming ambiguities - use table aliases

Scoping rules

- (Scope applies inwards) Table alias decalred in subquery Q can only be used in Q, or subqueries nested within Q
- If same table alias is declared in subquery Q and in an outer query, then use the declaration in Q

Scalar subqueries

- Occurs when subquery returns a single value (1 row, 1 column)
- Can be used as expression in queries

Row constructors

- Allow subqueries to return more than one attribute
- Number of attributes in row constructor must match the one of the subquery

```
SELECT name, population, gdp FROM countries
WHERE ROW(population, gdp) > ANY (SELECT population, gdp
                                FROM countries
                                WHERE name in ('Germany',
                                                'France'));
```

- If comparison `op` is `=` or `<>`, then tuple comparison works as expected
- However, for `<`, `<=`, `>`, `>=`, the row elements are compared left-to-right, stopping as soon as an unequal or null pair of elements is found
 - If either of this pair of elements is null, the result of the row comparison is unknown (null)
 - Otherwise comparison of this pair of elements determines the result
- Can use `SELECT ROW(...)``op ROW(...)``as res` to test the functionality

Equivalent subqueries

- `expr IN (subquery) ≡ expr = ANY (subquery)`
- `expr1 op ANY(SELECT expr2 FROM ... WHERE ...)` `≡ EXISTS(SELECT * FROM ... WHERE ... AND expr1 op expr2)`

Sorting

- `ORDER BY attribute ASC (DESC)`
- `ORDER BY attr1 ASC, attr2 DESC` - 2nd sorting criteria only affects result if 1st sorting criteria has ambiguity

Ranking

- `LIMIT k`: return first *k* tuples of the result table
- `OFFSET k`: specify the position of the first tuple to be considered
- Used with `ORDER BY`
- `ORDER BY ... OFFSET 5 LIMIT 5` - get 6th to 10th tuples in sorting criteria

SQL (AGGREGATION)

- Computes a single value, given a set of tuples
- e.g. `MIN()`, `MAX()`, `AVG()`, `COUNT()`, `SUM()`
- `MIN(A)` does not consider null values for attribute *A*

Signatures

- `MIN`, `MAX` defined for all data types, returns same type as input
- `SUM` defined for all numeric data types
 - `SUM(INTEGER)-> BIGINT` | `SUM-REAL)-> REAL`
- `COUNT` defined for all data types; returns `BIGINT`

Handling NULL

Let *R* be a non-empty relation with attribute *A*.

Query	Interpretation	Result
SELECT MIN(A) FROM R;	Minimum non-null value in A	0
SELECT MAX(A) FROM R;	Maximum non-null value in A	42
SELECT AVG(A) FROM R;	Average of non-null values in A	12
SELECT SUM(A) FROM R;	Sum of non-null values in A	48
SELECT COUNT(A) FROM R;	Count of non-null values in A	4
SELECT COUNT(*) FROM R;	Count of rows in R	5
SELECT AVG(DISTINCT A) FROM R;	Average of distinct non-null values in A	15
SELECT SUM(DISTINCT A) FROM R;	Sum of distinct non-null values in A	45
SELECT COUNT(DISTINCT A) FROM R;	Count of distinct non-null values in A	3

Let *R, S* be two relations with an attribute *A*,

- where *R* is an empty relation, and
- *S* is a non-empty relation with *n* tuples, only null values for *A*

Query	Result	Query	Result
SELECT MIN(A) FROM R;	null	SELECT MIN(A) FROM S;	null
SELECT MAX(A) FROM R;	null	SELECT MAX(A) FROM S;	null
SELECT AVG(A) FROM R;	null	SELECT AVG(A) FROM S;	null
SELECT SUM(A) FROM R;	null	SELECT SUM(A) FROM S;	null
SELECT COUNT(A) FROM R;	0	SELECT COUNT(A) FROM S;	0
SELECT COUNT(*) FROM R;	0	SELECT COUNT(*) FROM S;	<i>n</i>

Grouping

- Logical partition of relation into groups, based on values for specified attributes
- Aggregate function applied over group, so 1 result tuple for each group
- Given `GROUP BY a1, a2, .. an`, 2 tuples *x* and *y* belong to the same group if for all *k* ∈ {1..*n*}, *x.ak* IS NOT DISTINCT FROM *y.ak* evaluates to TRUE
 - i.e. NULL is treated as a value
- If column *A* of table *R* appears in `SELECT` clause, one of the following conditions must hold:
 1. *A* appears in the `GROUP BY` clause
 2. *A* appears as input of an aggregation function in the `SELECT` clause
 3. Primary key of *R* appears in the `GROUP BY` clause

```
-- find lowest, highest, overall population size for each continent
SELECT continent,
       MIN(population) AS lowest,
       MAX(population) AS highest,
       SUM(population) AS overall
FROM countries
GROUP BY continent;
```

Having

- Conditions check for each group defined by **GROUP BY** clause
- Must be used with **GROUP BY**
- If column A of table R appears in **HAVING** clause, one of the following conditions must hold:
 - A appears in the **GROUP BY** clause
 - A appears as input of an aggregation function in the **HAVING** clause
 - Primary key of R appears in the **GROUP BY** clause

```
-- Find all routes served by >12 airlines
SELECT from_code, to_code, COUNT(*) AS num_airlines
FROM routes GROUP BY from_code, to_code HAVING COUNT(*)
> 12;
```

Conceptual evaluation of queries

FROM → WHERE → GROUP BY → HAVING → SELECT → ORDER BY → LIMIT/OFFSET

SQL (CONDITIONALS)

CASE expression

- Generic conditional, like if/else statements
- Used in SELECT, ORDER BY, etc

Regular if-else	Switch-like statement
<pre>CASE WHEN cond1 then res1 WHEN cond2 then res2 ... WHEN condN then resN ELSE res0 END</pre>	<pre>CASE expression WHEN val1 then res1 WHEN val2 then res2 ... WHEN valN then resN ELSE res0 END</pre>

COALESCE

- Returns first non-NULL value in list of input args
- Returns NULL if lit of input args are NULL

```
-- if NULL, then set type to other
SELECT type, COUNT(*) AS city_ount
FROM
  (SELECT COALESCE(type, 'other') AS type
   FROM cities) t
GROUP BY type;
```

NULLIF

- NULLIF**(v_1 , v_2) returns NULL if $v_1 = v_2$, otherwise returns v_1
- Common use case: convert special values (zero, empty string) to NULL values

```
-- Unknown values are represented as 0
-- so set them to NULL for aggregation to ignore
SELECT MIN(NULLIF(gini, 0)) AS min_gini,
       AVG(NULLIF(gini, 0)) AS avg_gini,
FROM countries;
```

SQL (STRUCTURING QUERIES)

Common Table Expressions (CTEs)

- Temporarily named query
- Multiple CTEs can be used within an SQL statement
- CTEs can refer previous CTEs, or can be not referred at all
- Improves readabiliity, debugging, maintenance

```
WITH IsolatedEuropeanCountries AS (
  SELECT n.iso2, n.name AS country
  FROM borders b, countries n
  WHERE b.country1_iso2 = n.iso2
        AND b.country2_iso2 IS NULL
        AND n.continent = 'Europe'),
     AirportsInIsolatedEuropeanCountries AS (
  SELECT n.country, c.name AS city, a.code, a.name AS airport
  FROM IsolatedEuropeanCountries n, cities c, airports a
  WHERE n.iso2 = c.country_iso2
        AND c.name = a.city
        AND c.country_iso2 = a.country_iso2),
     UnusedJustForFun AS (
  SELECT COUNT(*)
  FROM IsolatedEuropeanCountries)
SELECT i.country, i.city, i.airport
FROM AirportsInIsolatedEuropeanCountries i LEFT OUTER JOIN routes r
  ON i.code = r.to_code
WHERE r.to_code IS NULL;
```

- Each C_i is the name of a temporary table defined by Q_i
- Each C_i can reference any other C_j that has been declared before C_i
- SQL statement S can reference any possible subset of all C_i

Views

- Permanently named query (virtual table)
 - Query is stored (not result), and re-executed each time view is used
- Can be used like normal tables
 - No restriction when used in **SELECT** statements
 - Restrictions for **INSERT**, **UPDATE**, **DELETE** (updatable view)

Updatable view requirements

Satisfy all of

- Must have exactly 1 entry in its **FROM** list, which must be a table, or another updatable view
- Can only update one attribute of a particular row at a time
- No **WITH**, **DISTINCT**, **GROUP BY**, **HAVING**, **LIMIT**, **OFFSET**
- No set operations **UNION**, **INTERSECT**, **EXCEPT**
- No aggregate functions

- No constraint violations

```
CREATE VIEW view_name AS
SELECT .. FROM .. WHERE .. GROUP BY ..;
```

Universal quantification

- No support for universal quantification (e.g. find names of all users that have visited all countries)
- Transform query using logical equivalences
 - There does not exist a country that the user has not visited
- Alternative interpretation
 - Number of tuples in Visited for that user must match total number of countries

Logical equivalences

- Can be used for other types of queries

- e.g. $p \implies q \equiv \sim p \vee q$

Recursive queries

- Using CTEs

Find all airports that can be reached from SIN with 0..2 stops.
(limitation to max. 2 stops purely for performance reasons)

```
WITH RECURSIVE flight_path AS (
  SELECT from_code, to_code, 0 AS stops
  FROM connections
  WHERE from_code = 'SIN'
  UNION ALL
  SELECT c.from_code, c.to_code, p.stops+1
  FROM flight_path p, connections c
  WHERE p.to_code = c.from_code
  AND p.stops < 2
)
SELECT DISTINCT to_code, stops
FROM flight_path
ORDER BY stops ASC;
```

to_code	stops
PEK	0
BKK	0
FRA	0
...	...
KUA	0
DUB	1
PEK	1
SIN	1
...	...
MME	1
AMS	2
BKK	2
PER	2
...
ZYL	2

90 tuples

825 tuples

1,561 tuples